


CS273A: Machine Learning & Data Mining



Prof. Alexander Ihler
Fall 2024

UCI Donald Bren School of
Information & Computer Sciences

6TH ANNUAL INDUSTRY SHOWCASE OCTOBER 8-9 2024

Register for the 6th Annual Industry Showcase

[RSVP Today!](#)

- **Masters** Student Reception: Tues 10/8 @ 5:00-6:30pm (*RSVP!*)
- **PhD** Student Reception: Wed 10/9 @ 5:00-6:30pm (*RSVP!*)

Outline

How does ML work?

Ex: Centroid Classifier

Optimal Decisions (in theory)

Bayes Classifiers

Types of Errors

Outline

How does ML work?

Ex: Centroid Classifier

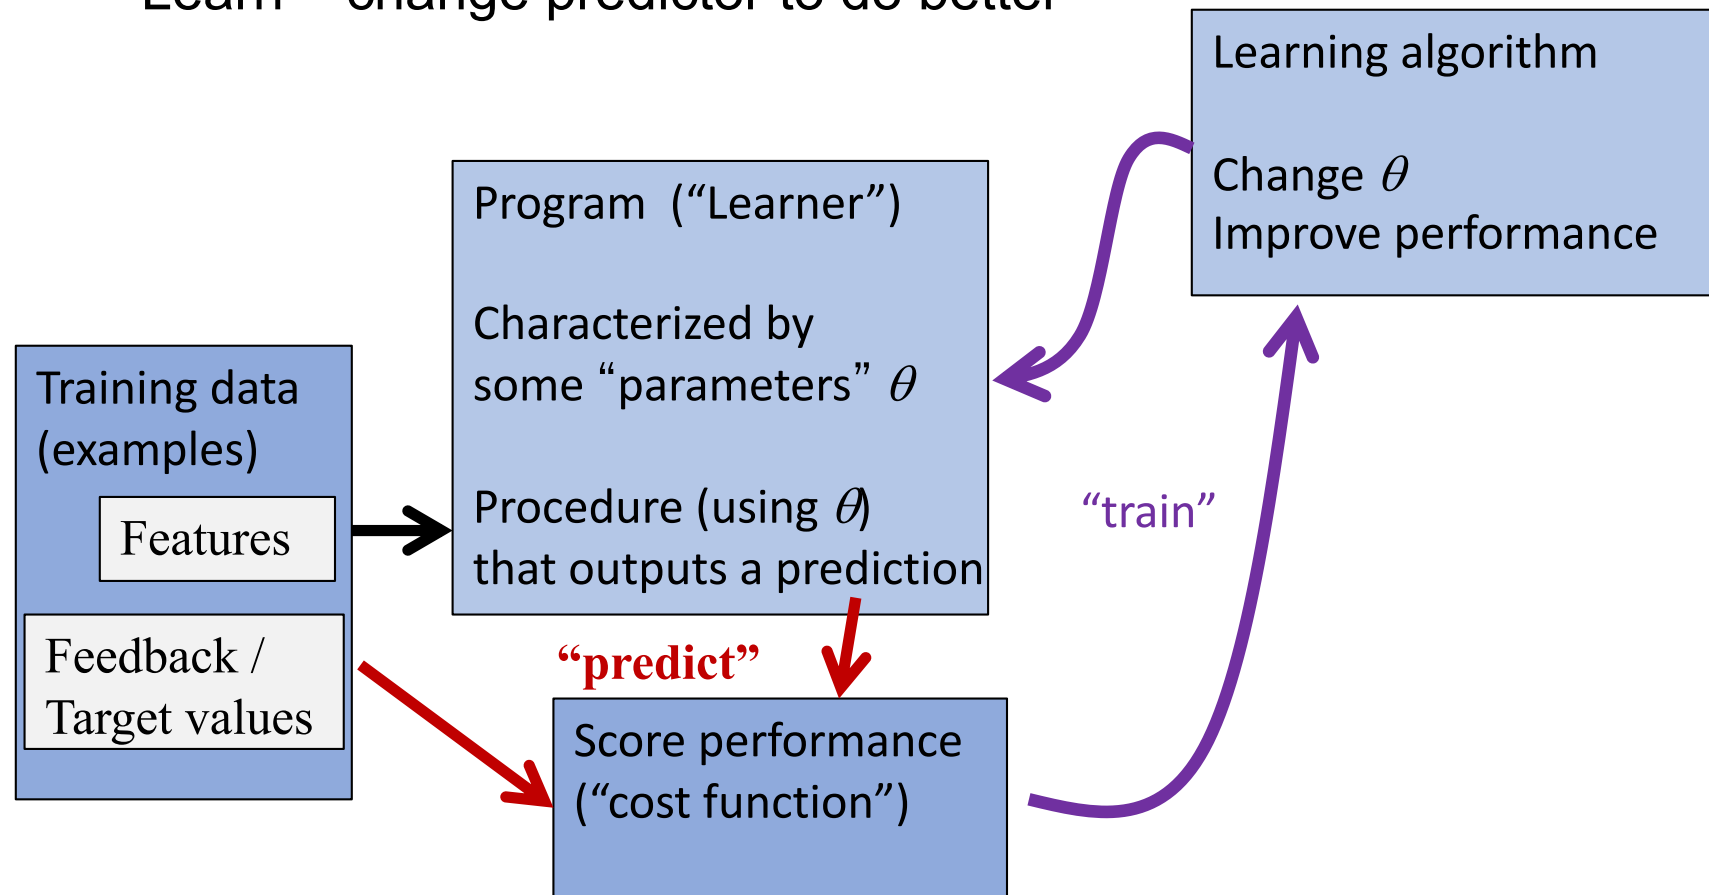
Optimal Decisions (in theory)

Bayes Classifiers

Types of Errors

How does machine learning work?

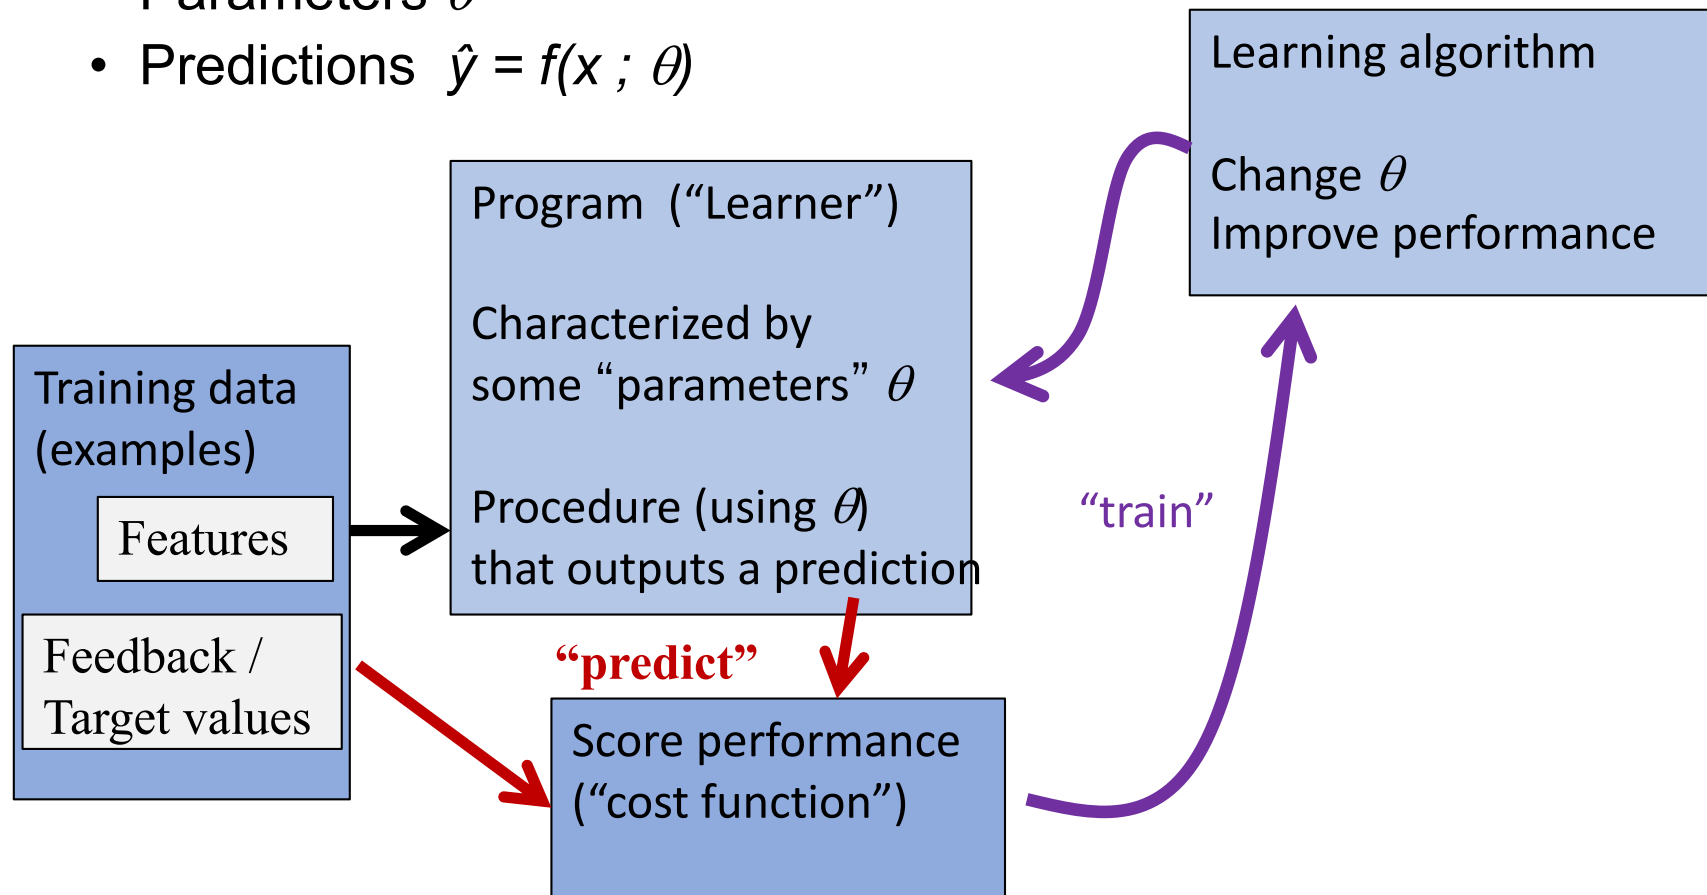
- “Meta-programming”
 - Predict – apply rules to examples
 - Score – get feedback on performance
 - Learn – change predictor to do better




Supervised Learning

- Notation

- Features x
- Targets y
- Parameters θ
- Predictions $\hat{y} = f(x ; \theta)$




Regression: scatter plots




- Data suggest a relationship between x and y
- *Prediction*: new x , what is y ?

Regression: nearest neighbor



- Find training datum $x^{(i)}$ closest to $x^{(new)}$; predict $y^{(i)}$

Regression: nearest neighbor



“Predictor”:


Given new features:
Find nearest example
Return its value

Parameters? Saved examples

Train on data X? Just save X

- Defines a function $f(x)$ implicitly
- “Form” is piecewise constant

Regression: linear regression



“Predictor”:

Evaluate line:

$$r = \theta_0 + \theta_1 x_1$$


return r

Parameters? Slope, intercept

Train on data X? Find a “close” l


- Define form of function $f(x)$ explicitly
- Find a good $f(x)$ within that family

Measuring error



Regression vs Classification

Regression




Features x

Real-valued target y

Predict continuous function $\hat{y}(x)$

Classification



Features x

Discrete class c
(usually 0/1 or +1/-1)

Predict discrete function $\hat{y}(x)$

Feature Vectors

$$x = (x_1, x_2, \dots, x_n)$$



A component of the vector, corresponding to the value of “feature 2”



n = dimensionality of the vector

Example 1:

Feature vector for a medical patient: $\mathbf{x} = (21.4, 6.1, 200)$
age height weight

Example 2:


Feature vector for a loan applicant: $\mathbf{x} = (21.4, 92697, 65k, 7.5k)$
age zipcode income debt

Feature Vectors as “Data Points”

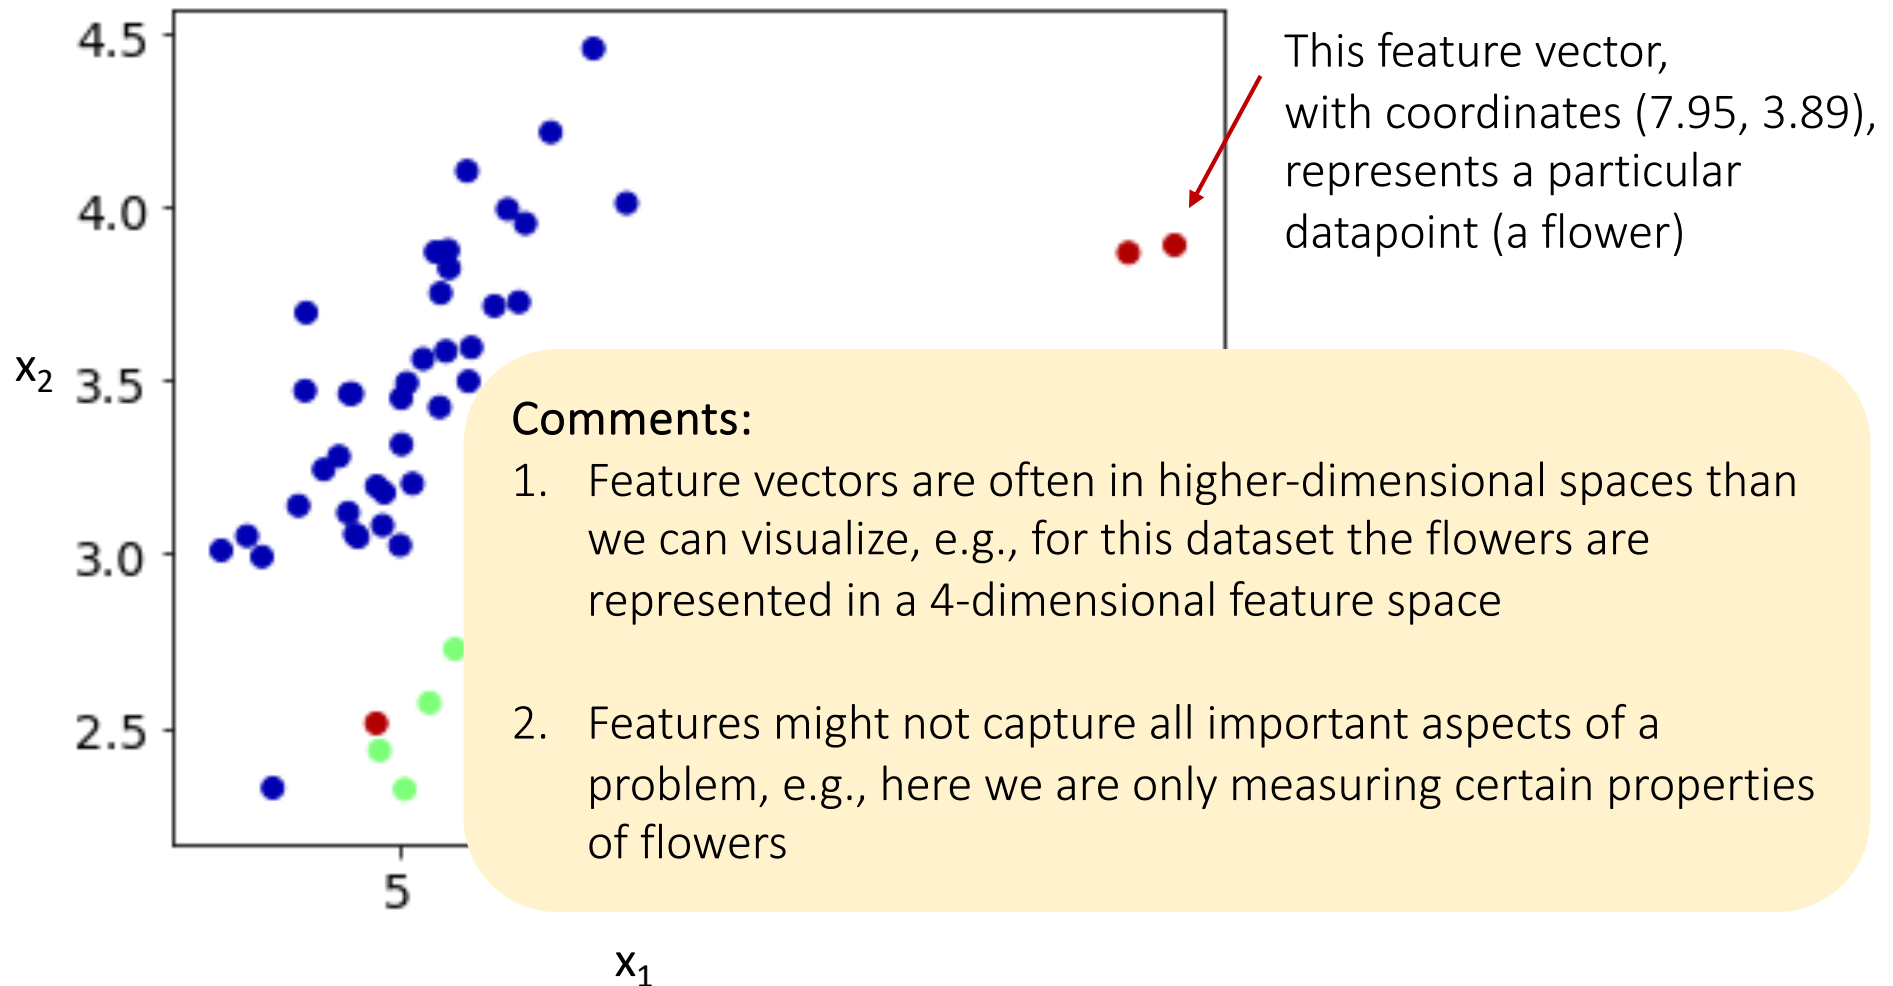
When we say “feature vector” we are referring to a point (in some d-dimensional space)

For example, if $n = 2$, we have $\mathbf{x} = (x_1, x_2)$


Here are two examples of vectors (red and blue)
representing two different datapoints in this 2-dimensional space



2-Dim Feature Space for Flowers




Classification



Classification

$$\text{ERR} = \frac{1}{m} \sum_i [y^{(i)} \neq \hat{y}(x^{(i)})]$$



Outline

How does ML work?

Ex: Centroid Classifier


Optimal Decisions (in theory)

Bayes Classifiers

Types of Errors

Ex: Centroid Classifier

- A simple, classical predictor
 - Train: decide what a “typical example” of each class y looks like
 - Identify the possible classes
 - For each class: “typical example” = the centroid (average) of those examples
 - Predict: which does the test point x look most like?
 - “most like” = closest in Euclidean distance



Mean Values of Individual Features

Visually, to compute the mean value of feature 1
.... we sum all values in 1st col and divide by m

$$\underline{X} = \begin{bmatrix} x_1^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

Mathematically,
mean value for feature j (e.g., j=2)

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

Sum over each row of
matrix **X** (index i is for rows)

Only sum the elements
of column j in matrix **X**

Means and Mean Vectors

Mean value for feature j

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

average over rows for column j (feature j) in X

Mean feature vector

(will also be referred to as the "centroid")

$$\mu = (\mu_1, \dots, \mu_n)$$

$$= \frac{1}{m} \sum_{i=1}^m x^{(i)}$$

average of the n vectors in X

Simple numerical example

$$X = \begin{pmatrix} x^{(1)} \\ x^{(2)} \end{pmatrix} = \begin{pmatrix} 21.4 & 6.1 & 200 \\ 28.1 & 5.5 & 145 \end{pmatrix}$$

Mean value for feature j = 1

$$\mu_1 = (21.4 + 28.1)/2 = 24.75$$

Mean feature vector (or centroid)

$$\mu = (24.75 \quad 5.8 \quad 172.5)$$


(in Python can use `np.mean(X,axis=0)`,
where X is a 2 x 3 array)

Ex: Centroid Classifier

- A simple, classical predictor


Training (“fit”):

```
train(X, y):  
  for each possible class c:  
    identify index of data points with y==c  
    compute centroid (mean)  $\mu_c$  of those data
```



Prediction:

```
predict(X):  # no known label y!  
  for data point x:  
    for each possible class c:  
      find distance of x to  $\mu_c$   
    pick the class c with smallest distance
```



Outline

How does ML work?

Ex: Centroid Classifier

Optimal Decisions (in theory)

Bayes Classifiers

Types of Errors

A simple, optimal classifier

- Classifier $f(x; \theta)$
 - maps observations x to predicted target values
- Simple example
 - Discrete feature x : $f(x; \theta)$ is a contingency table
 - Ex: spam filtering: observe just X_1 = sender in contact list?
- Suppose we knew the true conditional probabilities:
- Best prediction is the most likely target!

“Bayes error rate”

$$\begin{aligned} & \Pr[X=0] * \Pr[\text{wrong} \mid X=0] + \Pr[X=1] * \Pr[\text{wrong} \mid X=1] \\ &= \Pr[X=0] * (1 - \Pr[Y=S \mid X=0]) + \Pr[X=1] * (1 - \Pr[Y=K \mid X=1]) \end{aligned}$$


Can't do better than this without more information:
e.g., more features (email header, body text, etc.)

Feature	spam	keep
X=0	0.6	0.4
X=1	0.1	0.9

A simple classifier from data

- Training data $D=\{x^{(i)}, y^{(i)}\}$, Classifier $f(x; D)$
 - Discrete feature vector x
 - $f(x; D)$ is a contingency table
- Ex: Fisher Iris data, one feature
 - X_1 = sepal length (different ranges)
 - How should we make our predictions?
 - One method: just estimate the probabilities?

Sepal length	Iris setosa	Iris versicolor	Iris virginica
$X < 5$	21	30	5
$5 < X < 6$	23	21	30
$6 < X < 7$	0	16	35
$7 < X$	0	1	10



Sepal length	Iris setosa	Iris versicolor	Iris virginica
$X < 5$	0.375	0.536	0.089
$5 < X < 6$	0.311	0.284	0.405
$6 < X < 7$	0.	0.314	0.686
$7 < X$	0.	0.091	0.909

Estimating $p(y|X=x)$: “probabilistic” learning

(empirically estimated)

Gives a prediction *and* an (estimated) notion of confidence in that prediction

A simple classifier from data

- Training data $D=\{x^{(i)}, y^{(i)}\}$, Classifier $f(x; D)$
 - Discrete feature vector x
 - $f(x; D)$ is a contingency table
- Ex: Fisher Iris data, one feature
 - What if we bin the data more finely?
 - Find data falling within each range:

Two sources of error!


- Bayes error rate
(improve with more info in X)
- Mis-estimating probability
(improve with more data)

Sepal length	Iris setosa	Iris versicolor	Iris virginica
...			
5.25	0.57	0.07	0.36
5.5	0.09	0.48	0.43
5.75	0.08	0.38	0.54
...			

Sepal length	Iris setosa	Iris versicolor	Iris virginica
...			
5.48	1	0	0
5.5	0	0	1
5.52	0	0	0
...			


Inductive bias

- Allow us to extend observed data to unobserved ones
 - Interpolation / extrapolation
- What relationships do we expect in the data?
 - A (perhaps *the*) key question in ML models
 - Usually, data pull us away from assumptions only with evidence!




All of these explain the data in some way!


Overfitting & Complexity




Overfitting & Complexity




Overfitting & Complexity




Overfitting & Complexity



Overfitting & Complexity




How Overfitting Affects Prediction







Recall: Inductive bias

- How can we transfer observations to other, unobserved values?




Data



Models that “explain” data, extending to other values



- For $p(x,y)$? One option: discretize (histograms)



- Binning “transfers” data density to nearby feature values
- Too few bins = lose information; too many = noisy, no estimates at many locations

Fundamental issue of ML: How can we transfer information from “similar” examples?

Outline

How does ML work?

Ex: Centroid Classifier

Optimal Decisions (in theory)

Bayes Classifiers


Types of Errors

Gaussian probability models

- Estimate parameters of a Gaussian distribution from data
 - Gaussian dist: $\mathcal{N}(x; \mu_c, \sigma_c^2) = \left(2\pi\sigma_c^2\right)^{-\frac{1}{2}} \exp\left[-\frac{1}{2}(x - \mu_c)^2/\sigma_c^2\right]$
 - Empirical (and maximum likelihood) parameter estimates:

$$\hat{p}(Y = 1) = \frac{m_1}{m} \quad \hat{\mu}_1 = \frac{1}{m_1} \sum_{i:y^{(i)}=1} x^{(i)} \quad \hat{\sigma}_1^2 = \frac{1}{m_1} \sum_{i:y^{(i)}=1} (x^{(i)} - \hat{\mu}_1)^2$$

(and similarly for class 0)



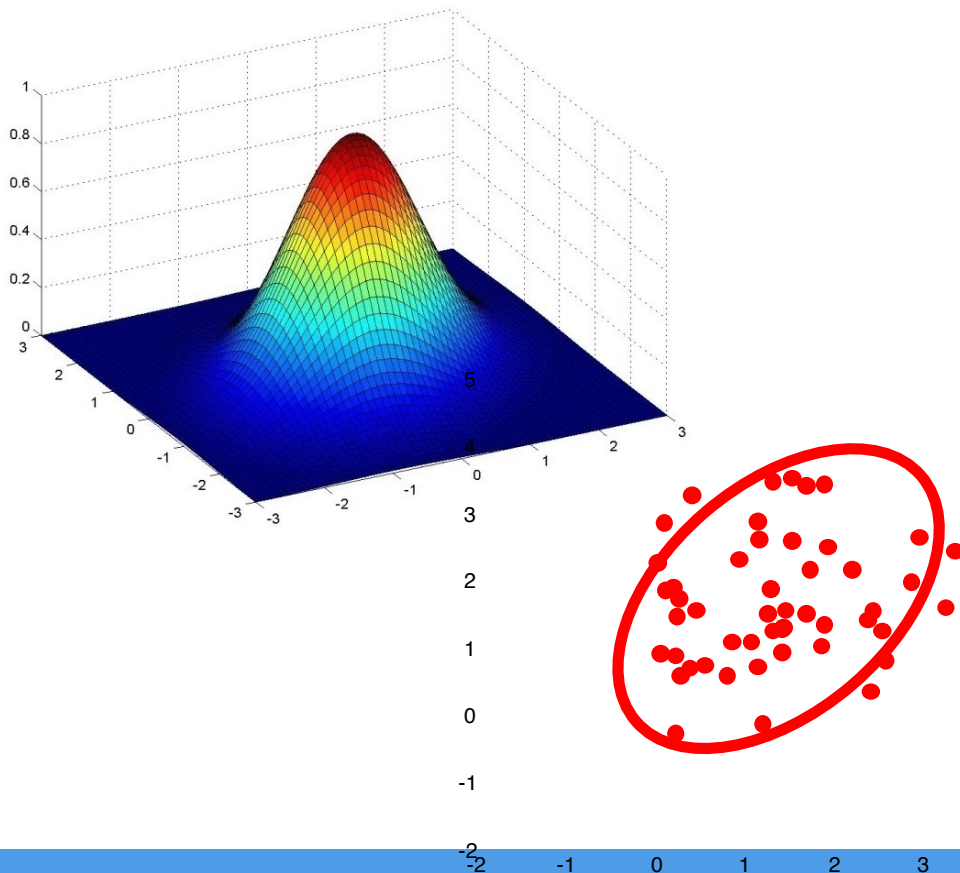
Multivariate Gaussian models

- Similar to univariate case

$$\mathcal{N}(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2}} |\Sigma|^{-\frac{1}{2}} \exp \left\{ -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right\}$$

$\mu = n \times 1$ mean vector

$\Sigma = n \times n$ covariance matrix



Maximum likelihood estimate:

$$\hat{\mu} = \frac{1}{m} \sum_j x^{(j)}$$

$$\hat{\Sigma} = \frac{1}{m} \sum_j (x^{(j)} - \hat{\mu})^T (x^{(j)} - \hat{\mu})$$

Bayes rule

- How to compute the probability of a hidden “cause” Y , after observing some evidence “effect” X :

$$p(Y|X) p(X) = p(X, Y) = p(X|Y) p(Y)$$

How probable is the hidden cause?

How often does Y cause X ?

$$\Rightarrow p(Y|X) = \frac{p(X|Y) p(Y)}{p(X)}$$

“Bayes rule”

- Example: flu
 - $P(F)$, $P(H|F)$
 - $P(F=1 | H=1) = ?$

$$= \frac{0.50 * 0.05}{0.50 * 0.05 + 0.20 * 0.95} = 0.116$$

F	P(F)
0	0.95
1	0.05

F	H	P(H F)
0	0	0.80
0	1	0.20
1	0	0.50
1	1	0.50

Bayes Classifiers from Data

- Estimate prior probability of each class, $p(y)$
 - E.g., how common is each type of Iris?
- Distribution of features given the class, $p(x | y=c)$
 - How likely are we to see “x” in each type of iris?
- Joint distribution $p(y|x)p(x) = p(x, y) = p(x|y)p(y)$
- Bayes Rule: $\Rightarrow p(y|x) = p(x|y)p(y)/p(x)$

(Use the rule of total probability to calculate the denominator!) \longrightarrow

$$= \frac{p(x|y)p(y)}{\sum_c p(x|y=c)p(y=c)}$$

Example: Gaussian Bayes, Iris Data


- Fit Gaussian distribution to each class {0,1,2}

$$p(y) = \text{Discrete}(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$$

$$p(x_1, x_2 | y = 0) = \mathcal{N}(x; \mu_0, \Sigma_0)$$

$$p(x_1, x_2 | y = 1) = \mathcal{N}(x; \mu_1, \Sigma_1)$$

$$p(x_1, x_2 | y = 2) = \mathcal{N}(x; \mu_2, \Sigma_2)$$



Then, Bayes rule:

$$p(Y = b|x) = \frac{\overset{\text{(How well does } Y=\text{blue explain } x?)}}{p(Y = b)p(x|Y = b)}}{\underset{\text{(How well do } Y=\text{green or } Y=\text{red explain } x?)}}{p(Y = b)p(x|Y = b) + p(Y = g)p(x|Y = g) + p(Y = r)p(x|Y = r)}}$$

Homework: Centroid Classifier

- Simple, special case of Gaussian Bayes classifier
- Estimate just the mean (centroid) of each data class

- Then, rule is simply:


predict class y by:

$$\hat{y}(x) = \arg \min_c \|x - \mu_c\|^2$$

Typically, use Euclidean distance:

$$\|x - \mu\|^2 = \sum_j (x_j - \mu_j)^2$$

though other distances also possible (more later...)



What about discrete features?

- Estimate joint probability for each class
 - E.g., how many times (what fraction) did each outcome occur?
- m data $\ll 2^n$ parameters?
- What about the zeros?
 - We learn that certain combinations are impossible?
 - What if we see these later in test data?
- Overfitting!

A	B	C	$p(A,B,C \mid Y=1)$
0	0	0	4/10
0	0	1	1/10
0	1	0	0/10
0	1	1	0/10
1	0	0	1/10
1	0	1	2/10
1	1	0	1/10
1	1	1	1/10

What about discrete features?

- Estimate joint probability for each class

- E.g., how many times (what fraction) did each outcome occur?

A	B	C	p(A,B,C Y=1)
0	0	0	4/10
0	0	1	1/10
0	1	0	0/10
0	1	1	0/10
1	0	0	1/10
1	0	1	2/10
1	1	0	1/10
1	1	1	1/10

- m data $\ll 2^n$ parameters?

- What about the zeros?

- We learn that certain combinations are impossible?
 - What if we see these later in test data?

- One option: regularize $\hat{p}(a, b, c) \propto (M_{abc} + \alpha)$

- Normalize to make sure values sum to one...


Naïve Bayes Classifiers

- Another option: reduce the model complexity by assuming the features are (conditionally) independent of one another
- Independence: $p(a,b) = p(a) p(b)$
- $p(x_1, x_2, \dots x_N | y=1) = p(x_1 | y=1) p(x_2 | y=1) \dots p(x_N | y=1)$
- Only need to estimate each individually

A	$p(A Y=1)$
0	.4
1	.6

B	$p(B Y=1)$
0	.7
1	.3

C	$p(C Y=1)$
0	.1
1	.9



A	B	C	$p(A,B,C Y=1)$
0	0	0	$.4 * .7 * .1$
0	0	1	$.4 * .7 * .9$
0	1	0	$.4 * .3 * .1$
0	1	1	...
1	0	0	
1	0	1	
1	1	0	
1	1	1	

Example: Naïve Bayes

Observed Data:

x_1	x_2	y
1	1	0
1	0	0
1	0	1
0	0	0
0	1	1
1	1	0
0	0	1
1	0	1

$$\hat{p}(y = 1) = \frac{4}{8} = (1 - \hat{p}(y = 0))$$

$$\hat{p}(x_1, x_2 | y = 0) = \hat{p}(x_1 | y = 0) \hat{p}(x_2 | y = 0)$$

$$\hat{p}(x_1 = 1 | y = 0) = \frac{3}{4}$$

$$\hat{p}(x_1 = 1 | y = 1) = \frac{2}{4}$$

$$\hat{p}(x_2 = 1 | y = 0) = \frac{2}{4}$$

$$\hat{p}(x_2 = 1 | y = 1) = \frac{1}{4}$$

Prediction given some observation x ?

$$\hat{p}(y = 1) \hat{p}(x = 11 | y = 1) \quad \begin{matrix} < \\ > \end{matrix} \quad \hat{p}(y = 0) \hat{p}(x = 11 | y = 0)$$

$$\frac{4}{8} \times \frac{2}{4} \times \frac{1}{4} \quad \quad \quad \frac{4}{8} \times \frac{3}{4} \times \frac{2}{4}$$

Decide class 0

Example: Naïve Bayes

Observed Data:

x_1	x_2	y
1	1	0
1	0	0
1	0	1
0	0	0
0	1	1
1	1	0
0	0	1
1	0	1

$$\hat{p}(y = 1) = \frac{4}{8} = (1 - \hat{p}(y = 0))$$

$$\hat{p}(x_1, x_2 | y = 0) = \hat{p}(x_1 | y = 0) \hat{p}(x_2 | y = 0)$$

$$\hat{p}(x_1 = 1 | y = 0) = \frac{3}{4} \quad \hat{p}(x_1 = 1 | y = 1) = \frac{2}{4}$$

$$\hat{p}(x_2 = 1 | y = 0) = \frac{2}{4} \quad \hat{p}(x_2 = 1 | y = 1) = \frac{1}{4}$$

$$\begin{aligned} \hat{p}(y = 1 | x_1 = 1, x_2 = 1) &= \frac{\frac{4}{8} \times \frac{2}{4} \times \frac{1}{4}}{\frac{3}{4} \times \frac{2}{4} \times \frac{4}{8} + \frac{2}{4} \times \frac{1}{4} \times \frac{4}{8}} \\ &= \frac{1}{4} \end{aligned}$$

Example: Joint Bayes

Observed Data:

x_1	x_2	y
1	1	0
1	0	0
1	0	1
0	0	0
0	1	1
1	1	0
0	0	1
1	0	1

$$\hat{p}(y = 1) = \frac{4}{8} = (1 - \hat{p}(y = 0))$$

$$\hat{p}(x_1, x_2 | y = 0) =$$

x_1	x_2	$p(x y=0)$
0	0	1/4
0	1	0/4
1	0	1/4
1	1	2/4

$$\hat{p}(x_1, x_2 | y = 1) =$$

x_1	x_2	$p(x y=1)$
0	0	1/4
0	1	1/4
1	0	2/4
1	1	0/4

$$\hat{p}(y = 1 | x_1 = 1, x_2 = 1) = \frac{\frac{4}{8} \times 0}{\frac{2}{4} \times \frac{4}{8} + 0 \times \frac{4}{8}} = 0$$

Naïve Bayes Models

- Variable y to predict, e.g. “auto accident in next year?”
- *Many* co-observed variables $x=[x_1 \dots x_n]$
 - Age, income, education, zip code, ...
- Learn $p(y \mid x_1 \dots x_n)$, to predict y ?
 - Arbitrary distribution: $O(d^n)$ values!

- Naïve Bayes:

$$p(y|x) = \frac{p(x|y)p(y)}{p(x)}$$

Bayes Rule

$$p(x|y) = \prod_j p(x_j|y)$$

“Naïve” : conditional independence

Now only $2 \cdot n \cdot d$ parameters!

- Note: may not be a good model of the data
 - Doesn't capture correlations in features
 - Can't capture some dependencies
- But in practice it often does quite well!

Outline

How does ML work?

Ex: Centroid Classifier

Optimal Decisions (in theory)

Bayes Classifiers

Types of Errors


Bayes Classifiers

- Bayes classification decision rule compares probabilities:

$$p(y = 0|x) \begin{matrix} < \\ > \end{matrix} p(y = 1|x)$$

$$= p(y = 0, x) \begin{matrix} < \\ > \end{matrix} p(y = 1, x)$$

- Can visualize this nicely if x is a scalar:




Bayes Classifiers

- Not all errors are created equally...
- Risk associated with each outcome?

Add multiplier alpha:

$$\alpha \begin{cases} p(y = 0, x) < \\ > \end{cases} p(y = 1, x)$$



False positive rate: $(\# y=0, \hat{y}=1) / (\# y=0)$


False negative rate: $(\# y=1, \hat{y}=0) / (\# y=1)$

Bayes Classifiers

- Increase alpha: prefer class 0
- Spam detection

Add multiplier alpha:

$$\alpha \ p(y = 0, x) \begin{matrix} < \\ > \end{matrix} p(y = 1, x)$$



False positive rate: $(\# y=0, \hat{y}=1) / (\# y=0)$


False negative rate: $(\# y=1, \hat{y}=0) / (\# y=1)$

Bayes Classifiers

- Decrease alpha: prefer class 1
- Cancer detection

Add multiplier alpha:

$$\alpha \ p(y = 0, x) \begin{matrix} < \\ > \end{matrix} p(y = 1, x)$$



False positive rate: $(\# y=0, \hat{y}=1) / (\# y=0)$

False negative rate: $(\# y=1, \hat{y}=0) / (\# y=1)$

Measuring Errors

- Confusion matrix
- Can extend to more classes

	Predict 0	Predict 1
Y=0	380	5
Y=1	338	3

- True positive rate: $\#(y=1, \hat{y}=1) / \#(y=1)$ -- “sensitivity”
- False negative rate: $\#(y=1, \hat{y}=0) / \#(y=1)$
- False positive rate: $\#(y=0, \hat{y}=1) / \#(y=0)$
- True negative rate: $\#(y=0, \hat{y}=0) / \#(y=0)$ -- “specificity”

Likelihood Ratio Tests

- Connection to classical, statistical decision theory:


$$p(y = 0, x) \underset{>}{\overset{<}{\leq}} p(y = 1, x) \quad = \quad \log \frac{p(y = 0)}{p(y = 1)} \underset{>}{\overset{<}{\leq}} \log \frac{p(x|y = 1)}{p(x|y = 0)}$$

“log likelihood ratio”

- Likelihood ratio: relative support for observation “x” under “alternative hypothesis” $y=1$, compared to “null hypothesis” $y=0$
- Can vary the decision threshold: $\gamma \underset{>}{\overset{<}{\leq}} \log \frac{p(x|y = 1)}{p(x|y = 0)}$
- Classical testing:
 - Choose gamma so that FPR is fixed (“p-value”)
 - Given that $y=0$ is true, what’s the probability we decide $y=1$?


ROC Curves

- Characterize performance as we vary the decision threshold?




Types of Supervised Learning

Probabilistic
Generative Learning




Full “generative” model
Also explain features,
e.g., $p(y, x)$

Probabilistic
Discriminative Learning



“Soft” predictions
Probability / confidence,
e.g., $p(y|x)$

Discriminative Learning




“Hard” (discrete) predictions
Minimize loss, e.g., error rate

Confidence predictions allow us to change our desired loss “after” training:

- Care more about one type of error than another?
- Expect more of one class than the other?
- (Easier to) combine different predictions? (see: ensembles)

ROC Curves

- Characterize performance as we vary our confidence threshold?



Questions?