




CS273A: Linear Regression



Prof. Alexander Ihler
Fall 2024

Linear Regression

Linear Regression via Least Squares

Gradient Descent Algorithms

Direct Minimization of Squared Error

Regression with Non-linear Features

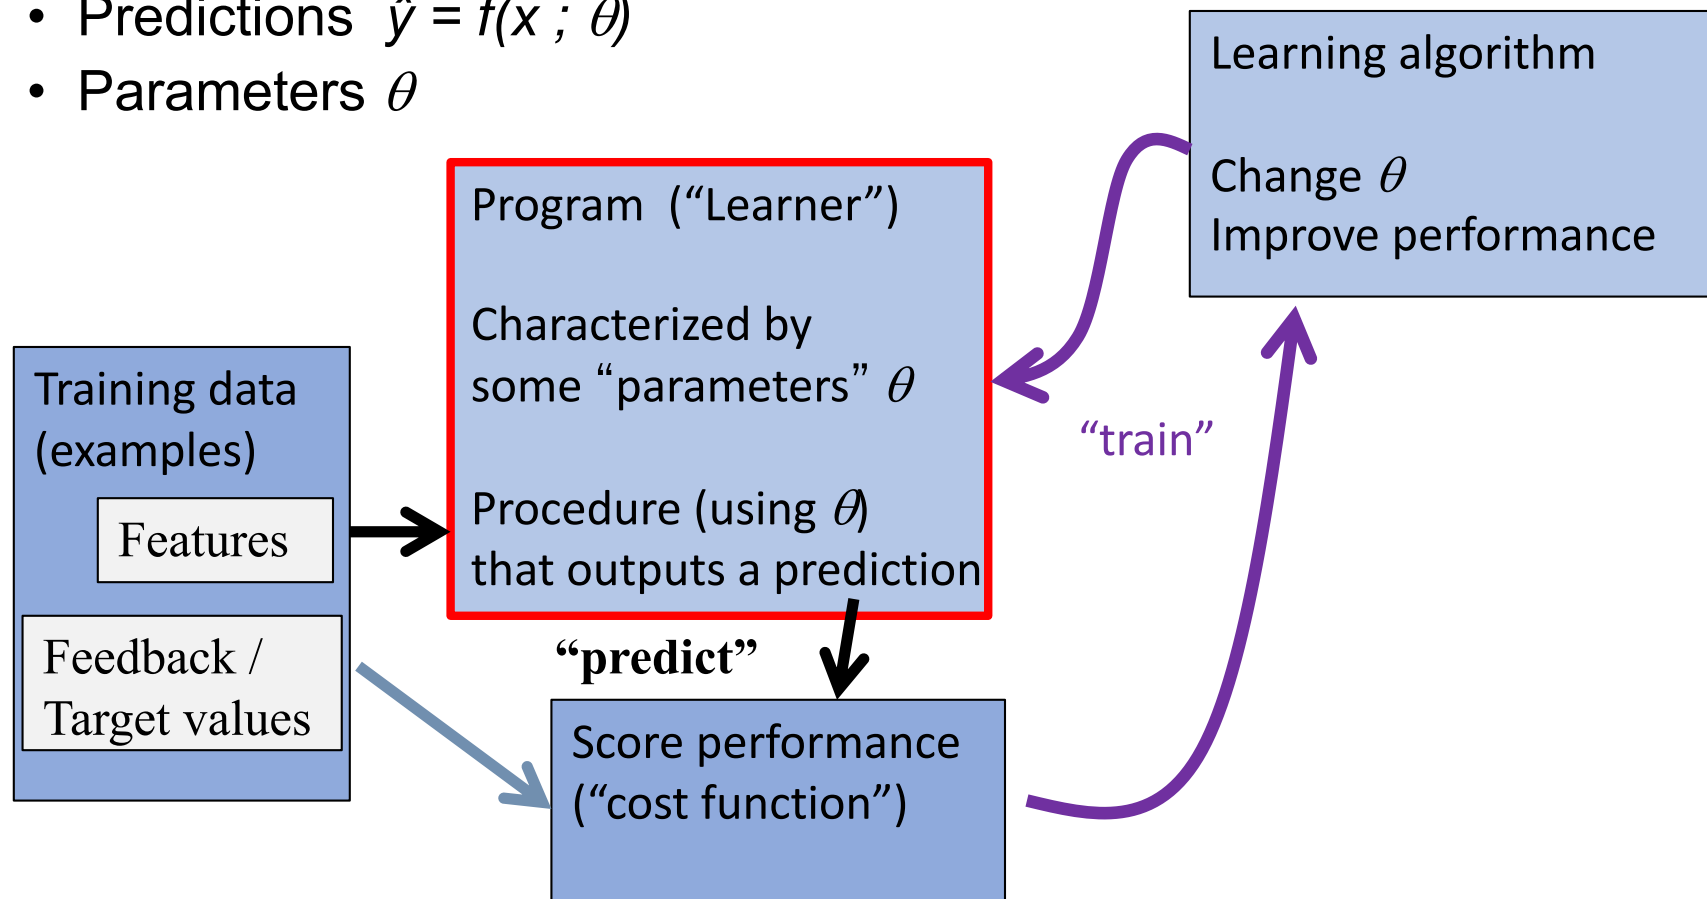
Bias, Variance, & Validation

Regularized Linear Regression


Supervised learning

- Notation

- Features x
- Targets y
- Predictions $\hat{y} = f(x; \theta)$
- Parameters θ



Linear regression



“Predictor”:



Evaluate line:

$$r = \theta_0 + \theta_1 x_1$$

return r

- Define form of function $f(x)$ explicitly
- Find a good $f(x)$ within that family

More dimensions?



“Predictor”:

Evaluate linear response:

$$r = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

return r

Notation

$$\hat{y}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots$$

Define “feature” $x_0 = 1$ (constant)

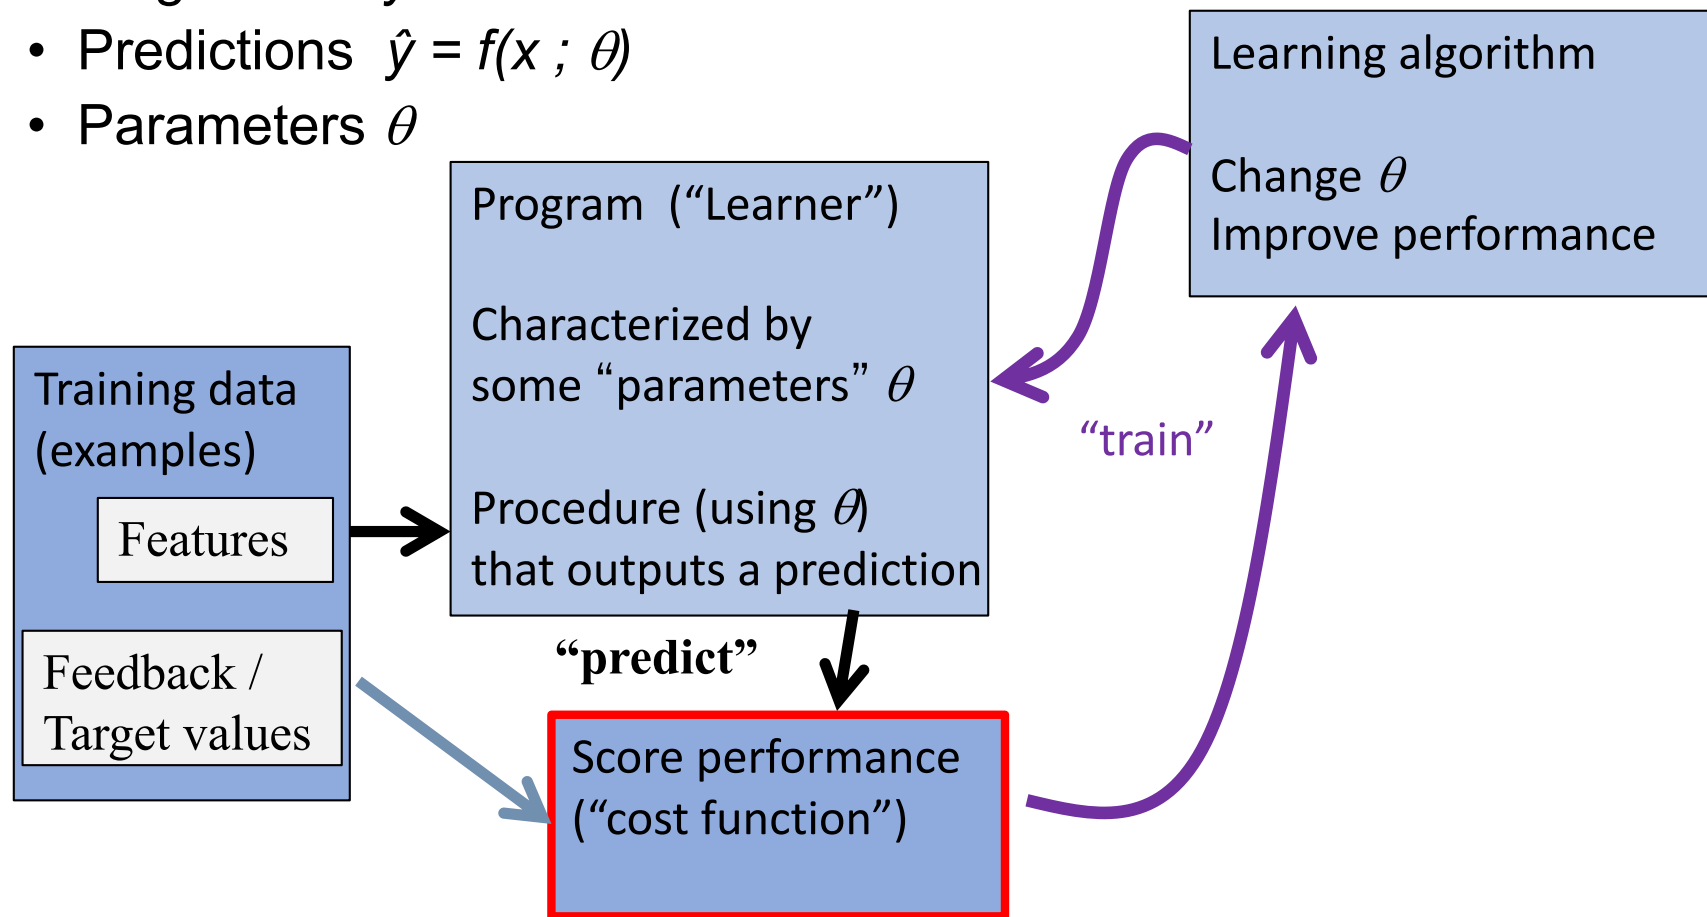
Then

$$\hat{y}(x) = \theta x^T$$
$$\underline{\theta} = [\theta_0, \dots, \theta_n]$$
$$\underline{x} = [1, x_1, \dots, x_n]$$


Supervised learning

- Notation

- Features x
- Targets y
- Predictions $\hat{y} = f(x; \theta)$
- Parameters θ



Measuring error



Mean squared error

- How can we quantify the error?

$$\begin{aligned}\text{MSE}, J(\underline{\theta}) &= \frac{1}{m} \sum_j (y^{(j)} - \hat{y}(x^{(j)}))^2 \\ &= \frac{1}{m} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)T})^2\end{aligned}$$

- Could choose something else, of course...
 - Computationally convenient (more later)
 - Measures the variance of the residuals
 - Corresponds to likelihood under Gaussian model of “noise”

$$\mathcal{N}(y ; \mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2} (y - \mu)^2 \right\}$$

MSE cost function

$$\begin{aligned}\text{MSE, } J(\underline{\theta}) &= \frac{1}{m} \sum_j (y^{(j)} - \hat{y}(x^{(j)}))^2 \\ &= \frac{1}{m} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)T})^2\end{aligned}$$

- Rewrite using matrix form

$$\underline{\theta} = [\theta_0, \dots, \theta_n]$$

$$\underline{y} = [y^{(1)} \dots, y^{(m)}]^T$$

$$\underline{X} = \begin{bmatrix} x_0^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_0^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

$$J(\underline{\theta}) = \frac{1}{m} (\underline{y}^T - \underline{\theta} \underline{X}^T) \cdot (\underline{y}^T - \underline{\theta} \underline{X}^T)^T$$

Python / NumPy:

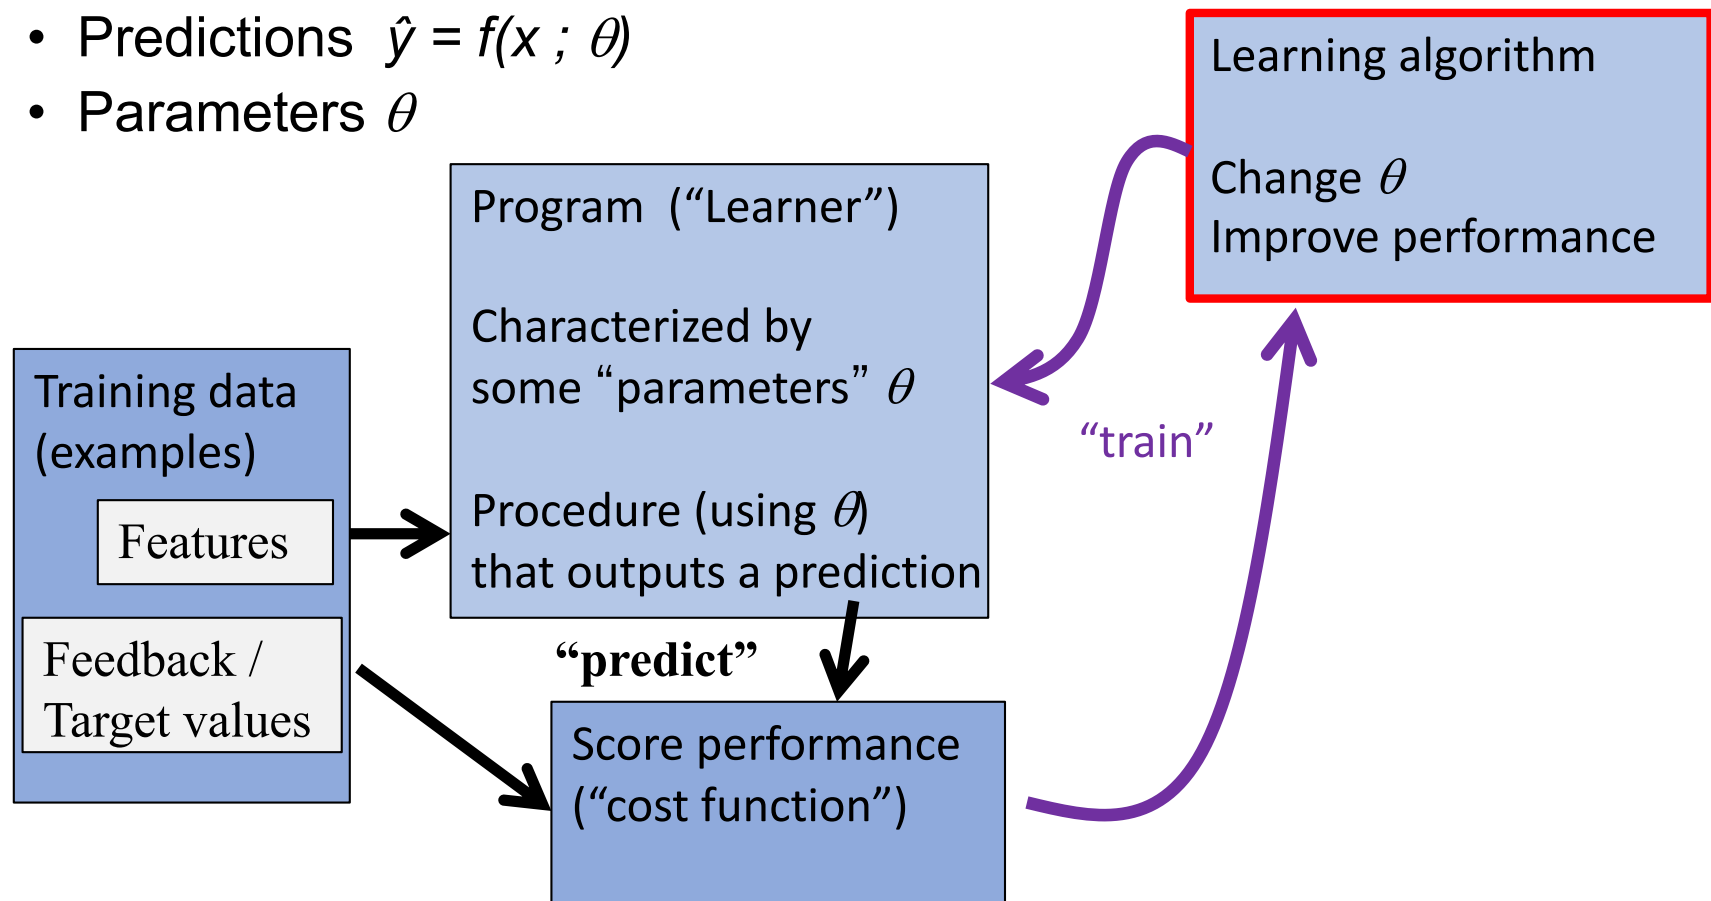
```
e = Y - X.dot( theta.T );
```

```
J = e.T.dot( e ) / m # = np.mean( e ** 2 )
```





Supervised learning

- Notation

- Features x
- Targets y
- Predictions $\hat{y} = f(x; \theta)$
- Parameters θ




Visualizing the cost function



Finding good parameters

- Want to find parameters which minimize our error...
- Think of a cost “surface”: error residual for that θ ...



Linear Regression

Linear Regression via Least Squares

Gradient Descent Algorithms


Direct Minimization of Squared Error

Regression with Non-linear Features

Bias, Variance, & Validation


Regularized Linear Regression

Gradient descent



- How to change θ to improve $J(\theta)$?
- Choose a direction in which $J(\theta)$ is decreasing

Gradient descent




- How to change θ to improve $J(\theta)$?
- Choose a direction in which $J(\theta)$ is decreasing
- Derivative $\frac{\partial J(\theta)}{\partial \theta}$
- Positive \Rightarrow increasing
- Negative \Rightarrow decreasing

Gradient descent in more dimensions

- Gradient vector

$$\nabla J(\underline{\theta}) = \left[\frac{\partial J(\underline{\theta})}{\partial \theta_0} \quad \frac{\partial J(\underline{\theta})}{\partial \theta_1} \quad \dots \right]$$



Indicates direction of steepest ascent
(negative = steepest descent)

Gradient descent


- Initialization
- Step size α
 - Can change over iterations
- Gradient direction
- Stopping condition

Initialize θ

Do{

$$\theta \leftarrow \theta - \alpha \nabla_{\theta} J(\theta)$$

} while $(\alpha \|\nabla_{\theta} J\| > \epsilon)$



Gradient for the MSE

- MSE $J(\underline{\theta}) = \frac{1}{m} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)})^2$

- $\nabla J = ?$ $J(\underline{\theta}) = \frac{1}{m} \sum_j \overbrace{(y^{(j)} - \theta_0 x_0^{(j)} - \theta_1 x_1^{(j)} - \dots)}^{e_j(\theta)}^2$

$$\frac{\partial J}{\partial \theta_0} = \frac{\partial}{\partial \theta_0} \frac{1}{m} \sum_j (e_j(\theta))^2$$

$$= \frac{1}{m} \sum_j \frac{\partial}{\partial \theta_0} (e_j(\theta))^2$$

$$= \frac{1}{m} \sum_j 2e_j(\theta) \frac{\partial}{\partial \theta_0} e_j(\theta)$$

$$\frac{\partial}{\partial \theta_0} e_j(\theta) = \cancel{\frac{\partial}{\partial \theta_0} y^{(j)}} - \frac{\partial}{\partial \theta_0} \theta_0 x_0^{(j)} - \cancel{\frac{\partial}{\partial \theta_0} \theta_1 x_1^{(j)}} - \dots$$

0 **0**

$$= -x_0^{(j)}$$

Gradient for the MSE

- MSE $J(\underline{\theta}) = \frac{1}{m} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)})^2$

- $\nabla J = ?$ $J(\underline{\theta}) = \frac{1}{m} \sum_j (y^{(j)} - \overbrace{\theta_0 x_0^{(j)} - \theta_1 x_1^{(j)} - \dots}^{e_j(\theta)})^2$

$$\begin{aligned} \nabla J(\underline{\theta}) &= \begin{bmatrix} \frac{\partial J}{\partial \theta_0} & \frac{\partial J}{\partial \theta_1} & \dots \end{bmatrix} \\ &= \begin{bmatrix} \underbrace{\frac{2}{m} \sum_j -e_j(\theta) x_0^{(j)}} & \underbrace{\frac{2}{m} \sum_j -e_j(\theta) x_1^{(j)}} & \dots \end{bmatrix} \end{aligned}$$

Gradient descent

- Initialization
- Step size α
 - Can change over iterations
- Gradient direction
- Stopping condition

Initialize θ

Do{

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J(\theta)$

} while $(\alpha \|\nabla_{\theta} J\| > \epsilon)$

$$J(\underline{\theta}) = \frac{1}{m} \sum_j (y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)T})^2$$

$$\nabla J(\underline{\theta}) = -\frac{2}{m} \sum_j \underbrace{(y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)T})}_{\text{Error magnitude \& direction for datum } j} \cdot \underbrace{[x_0^{(j)} \ x_1^{(j)} \ \dots]}_{\text{Sensitivity to each parameter}}$$

**Error magnitude &
direction for datum j**

**Sensitivity to
each parameter**

Derivative of MSE

- Rewrite using matrix form

$$\nabla J(\underline{\theta}) = -\frac{2}{m} \sum_j \underbrace{(y^{(j)} - \underline{\theta} \cdot \underline{x}^{(j)})^T}_{\text{Error magnitude \& direction for datum } j} \cdot \underbrace{[x_0^{(j)} \ x_1^{(j)} \ \dots]}_{\text{Sensitivity to each } \theta_i}$$

$$\underline{\theta} = [\theta_0, \dots, \theta_n]$$

$$\underline{y} = [y^{(1)} \ \dots \ y^{(m)}]^T$$



$$\nabla J(\underline{\theta}) = -\frac{2}{m} (\underline{y}^T - \underline{\theta} \underline{X}^T) \cdot \underline{X}$$

$$\underline{X} = \begin{bmatrix} x_0^{(1)} & \dots & x_n^{(1)} \\ \vdots & \ddots & \vdots \\ x_0^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$$

```
e = Y - X.dot( theta.T ) # error residual
DJ = - e.dot(X) * 2.0/m  # compute the gradient
theta -= alpha * DJ      # take a step
```



Gradient descent on cost function

Initialization





Gradient descent on cost function

Iteration 1





Gradient descent on cost function

Iteration 10





Gradient descent on cost function

Iteration 30




Gradient descent on cost function

Iteration 90






Comments on gradient descent

- Very general algorithm
 - We'll see it many times
- Local minima
 - Sensitive to starting point




Comments on gradient descent

- Very general algorithm
 - We'll see it many times
- Local minima
 - Sensitive to starting point
- Step size
 - Too large? Too small? Automatic ways to choose?
 - May want step size to decrease with iteration
 - Common choices:
 - Fixed
 - Linear: $C/(\text{iteration})$
 - Line search / backoff (Armijo, etc.)
 - Newton's method



Newton's method

- Want to find the roots of $f(x)$
 - “Root”: value of x for which $f(x)=0$




- Initialize to *some* point x
- Compute the tangent at x & compute where it crosses x -axis

$$\nabla f(z) = \frac{0 - f(z)}{z' - z} \Rightarrow z' = z - \frac{f(z)}{\nabla f(z)}$$

- Optimization: find roots of $\nabla J(\theta)$

$$\nabla \nabla J(\theta) = \frac{0 - \nabla J(\theta)}{\theta' - \theta} \Rightarrow \theta' = \theta - \frac{\nabla J(\theta)}{\nabla \nabla J(\theta)} \quad (\text{“Step size” } \lambda = 1/\nabla \nabla J ; \text{ inverse curvature})$$

- Does not always converge; sometimes unstable
- If converges, usually very fast
- Works well for smooth, non-pathological functions, locally quadratic
- For n large, may be computationally hard: $O(n^2)$ storage, $O(n^3)$ time



(Multivariate:

$\nabla J(\theta)$ = gradient vector

$\nabla^2 J(\theta)$ = matrix of 2nd derivatives

$a/b = a b^{-1}$, matrix inverse)

Stochastic / Online gradient descent

- MSE

$$J(\theta) = \frac{1}{m} \sum_i J^{(i)}(\theta) \quad J^{(i)}(\theta) = \left(y^{(i)} - \theta \cdot x^{(i)T} \right)^2$$

- Gradient

$$\nabla J(\theta) = \frac{1}{m} \sum_i \nabla J^{(i)}(\theta) \quad \nabla J^{(i)}(\theta) = \left(y^{(i)} - \theta \cdot x^{(i)T} \right) [x_0^{(i)} \ x_1^{(i)} \ \dots]$$

- Stochastic (or “online”) gradient descent:

- Use updates based on individual datum j , chosen at random
- At optima, $\mathbb{E}[\nabla J^{(i)}(\theta)] = \nabla J(\theta) = 0$
(average over the data)

Stochastic Gradient Descent

- Initialize theta
- Select a data point & update



Initialize θ

Do {

 for $i=1:m$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J^{(i)}(\theta)$

 } while (not done)



Stochastic Gradient Descent

- Initialize theta
- Select a data point & update,



Initialize θ

Do {

 for $i=1:m$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J^{(i)}(\theta)$

 } while (not done)



Stochastic Gradient Descent

- Initialize theta
- Select a data point & update,



Initialize θ

Do {

 for $i=1:m$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J^{(i)}(\theta)$

 } while (not done)



Stochastic Gradient Descent

- Initialize theta
- Select a data point & update,



Initialize θ

Do {

 for $i=1:m$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J^{(i)}(\theta)$

 } while (not done)



Stochastic Gradient Descent

- Initialize theta
- Select a data point & update,



Initialize θ

Do {

 for $i=1:m$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J^{(i)}(\theta)$

 } while (not done)



Stochastic Gradient Descent

$$J^{(i)}(\theta) = \left(y^{(i)} - \theta \cdot x^{(i)T} \right)^2$$

$$\nabla J^{(i)}(\theta) = 2 \left(y^{(i)} - \theta \cdot x^{(i)T} \right) \left[x_0^{(i)}, x_1^{(i)}, \dots \right]$$

Initialize θ

Do {

 for $i=1:m$

$\theta \leftarrow \theta - \alpha \nabla_{\theta} J^{(i)}(\theta)$

} while (not done)

- **Benefits**


- Lots of data = many more updates per pass
- Computationally faster

- **Drawbacks**


- No longer strictly “descent”
- Stopping conditions may be harder to evaluate
(Can use “running estimates” of $J(\cdot)$, etc.)

Mini-Batch Gradient Descent


- Select b data points at random & use average gradient
 - Interpolates between SGD & Batch GD
 - Fewer updates per epoch than SGD, but less noisy updates
- Example gradients (batch vs. **mini-batch**)



$b=1$
(SGD)



$b=3$



$b=5$

Linear Regression

Linear Regression via Least Squares

Gradient Descent Algorithms

Direct Minimization of Squared Error

Regression with Non-linear Features

Bias, Variance, & Validation


Regularized Linear Regression

MSE Minimum

- Consider a simple problem
 - One feature, two data points
 - Two unknowns: θ_0, θ_1
 - Two equations:

$$y^{(1)} = \theta_0 + \theta_1 x^{(1)}$$

$$y^{(2)} = \theta_0 + \theta_1 x^{(2)}$$



- Can solve this system directly:

$$\underline{y}^T = \underline{\theta} \underline{X}^T \quad \Rightarrow \quad \hat{\underline{\theta}} = y^T (\underline{X}^T)^{-1}$$

- However, most of the time, $m > n$
 - There may be no linear function that hits all the data exactly
 - Instead, solve directly for minimum of MSE function

MSE Minimum


$$\nabla J(\underline{\theta}) = -\frac{2}{m}(\underline{y}^T - \underline{\theta} \underline{X}^T) \cdot \underline{X} = \underline{0}$$

- Simplify with some algebra:

$$\underline{y}^T \underline{X} - \underline{\theta} \underline{X}^T \cdot \underline{X} = \underline{0}$$

$$\underline{y}^T \underline{X} = \underline{\theta} \underline{X}^T \cdot \underline{X}$$

$$\underline{\theta} = \underline{y}^T \underline{X} (\underline{X}^T \underline{X})^{-1}$$



- $\underline{X} (\underline{X}^T \underline{X})^{-1}$ is called the “pseudo-inverse”
- If \underline{X}^T is square and full rank, this is the inverse
- If $m > n$: overdetermined; gives minimum MSE fit
- $(\underline{X}^T \underline{X})$ not invertible? Underdetermined (multiple sol'ns)

Python MSE

- This is easy to solve in Python / NumPy...

$$\underline{\theta} = \underline{y}^T \underline{X} (\underline{X}^T \underline{X})^{-1}$$


```
# y = np.array( [[y1], ... , [ym]] )  
# X = np.array( [[x1_0 ... x1_n], [x2_0 ... x2_n], ...] )
```

```
# Solution 1: “manual”  
th = y.T @ X @ np.linalg.inv(X.T @ X)
```


```
# Solution 2: “least squares solve”  
th = np.linalg.lstsq(X, Y)
```

Effects of MSE choice

- Sensitivity to outliers



L1 error: Mean Absolute Error



Cost functions for regression


$$\ell_2 : (y - \hat{y})^2 \quad \text{(MSE)}$$

$$\ell_1 : |y - \hat{y}| \quad \text{(MAE)}$$

Something else entirely...

$$c - \log(\exp(-(y - \hat{y})^2) + c) \quad \text{(???)}$$

Arbitrary functions cannot be
solved in closed form
- use gradient descent



Linear Regression

Linear Regression via Least Squares

Gradient Descent Algorithms

Direct Minimization of Squared Error



Regression with Non-linear Features

Bias, Variance, & Validation

Regularized Linear Regression

Nonlinear functions

- What if our hypotheses are not lines?
 - Ex: higher-order polynomials



Nonlinear functions

- Single feature x , predict target y :

$$D = \{(x^{(j)}, y^{(j)})\}$$



Add features:

$$D = \{([x^{(j)}, (x^{(j)})^2, (x^{(j)})^3], y^{(j)})\}$$

$$\hat{y}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3$$



$$\hat{y}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

Linear regression in new features




- Sometimes useful to think of “feature transform”

$$\Phi(x) = [1, x, x^2, x^3, \dots]$$

$$\hat{y}(x) = \underline{\theta} \cdot \Phi(x)$$

Higher-order polynomials

- Fit in the same way
- More “features”







Features

- In general, can use any features we think are useful
- Other information about the problem
 - Anything you can encode as fixed-length vectors of numbers
- Polynomial functions
 - Features $[1, x, x^2, x^3, \dots]$
- Other functions
 - $1/x$, $\text{sqrt}(x)$, $x_1 * x_2, \dots$
- “Linear regression” = linear in the parameters
 - Features we can make as complex as we want!



Higher-order polynomials

- Are more features better?
- “Nested” hypotheses
 - 2nd order more general than 1st,
 - 3rd order more general than 2nd, ...
- Fits the observed data better





Overfitting and complexity

- More complex models will always fit the training data better
- But they may “overfit” the training data, learning complex relationships that are not really present




Test data

- After training the model
- Go out and get more data from the world
 - New observations (x,y)
- How well does our model perform?



Training versus test error

- Plot MSE as a function of model complexity
 - Polynomial order
- Decreases
 - More complex function fits training data better
- What about new data?
- 0th to 1st order
 - Error decreases
 - Underfitting
- Higher order
 - Error increases
 - Overfitting



Linear Regression

Linear Regression via Least Squares

Gradient Descent Algorithms

Direct Minimization of Squared Error



Regression with Non-linear Features

Bias, Variance, & Validation

Regularized Linear Regression


Inductive bias

- The assumptions needed to predict examples we haven't seen
- Makes us “prefer” one model over another
- Polynomial functions; smooth functions; etc
- Some bias is necessary for learning!

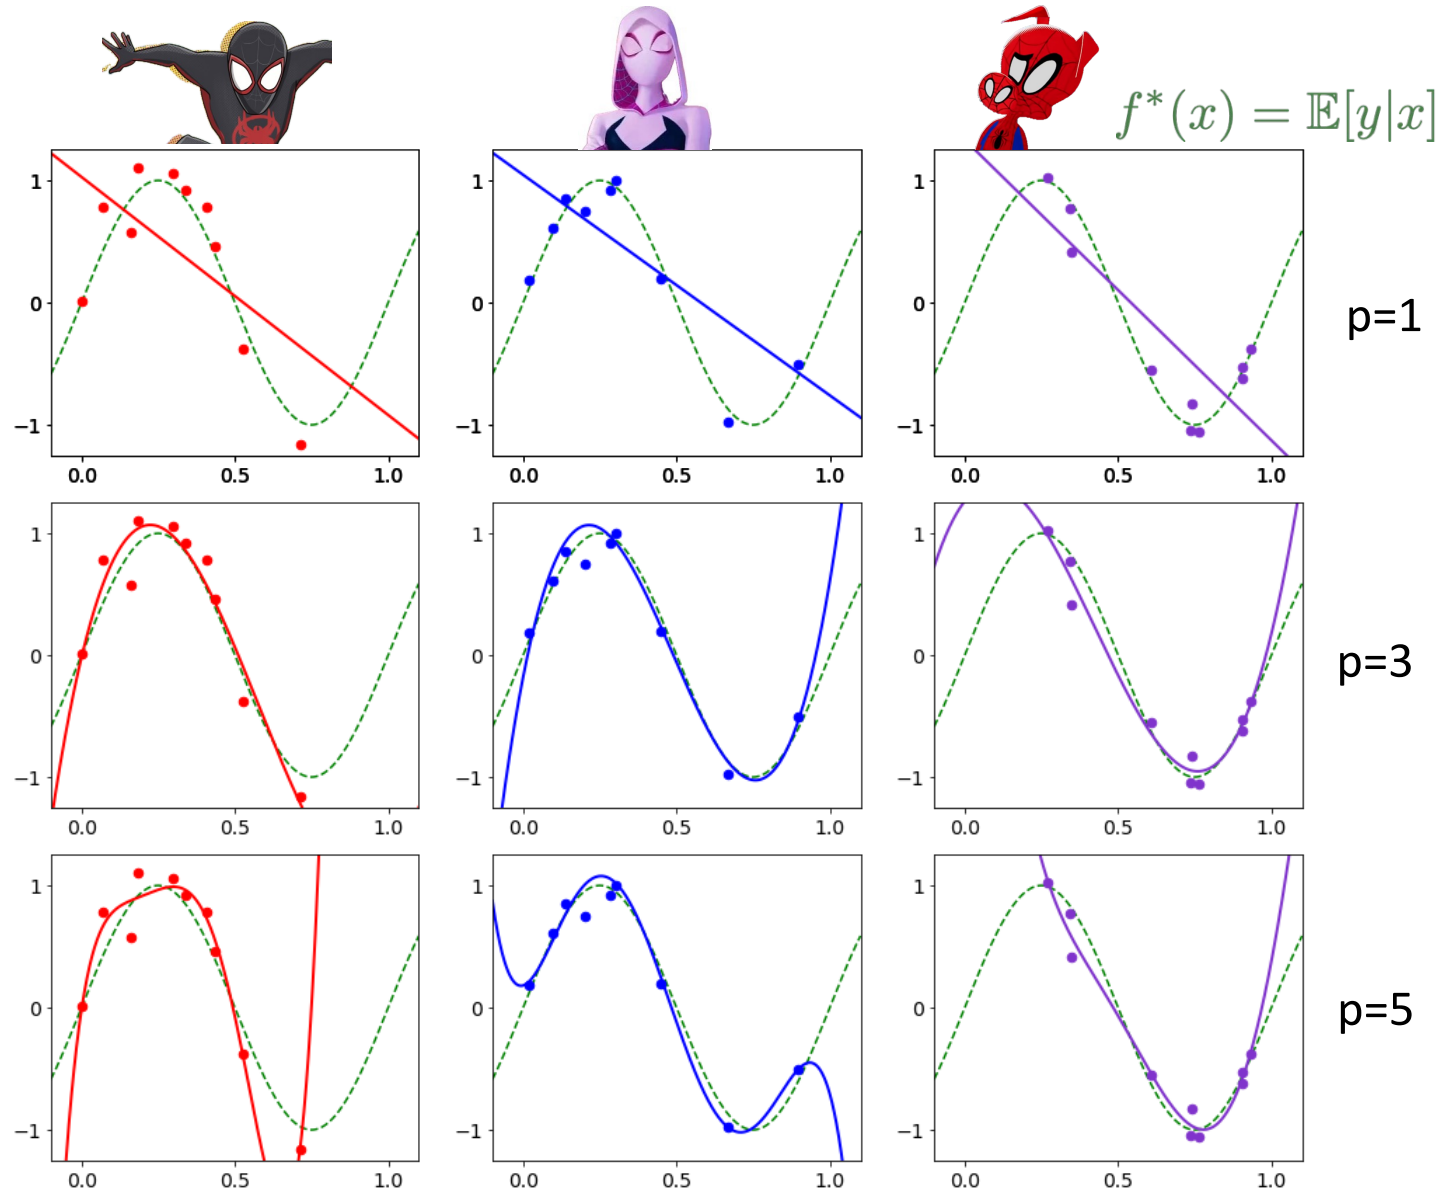


Bias & variance

We collect three different possible data sets:




$$\hat{y}(x) = \hat{\theta}_0 + \hat{\theta}_1 x$$




Bias & variance

- Two different types of “errors”



Low bias, high variance



High bias, low variance

- Same MSE (average distance to the bullseye) in both!

$$\frac{1}{m} \sum_i \|x^{(i)} - x^*\|^2 = \frac{1}{m} \sum_i \|\bar{x} - x^*\|^2 + \frac{1}{m} \sum_i \|x^{(i)} - \bar{x}\|^2$$

MSE


(Bias)²

Variance

Systematic error,
“on average”

Variation from
run to run


Bias & variance




Degree = 1

M = 10 data points


All models are similar to each other
Models are not very close to f^*




Bias & variance



Degree = 3
m = 10 data points




Bias & variance




Degree = 6
 $m = 10$ data points

Models are, **on average**, close to f^*
Models are very different from each other




Bias & variance



Degree = 6
m = 50 data points

Bias & variance are functions of the learner, training procedure, and training data size

Here: given more data per model, the models are far more similar (lower variance), while still having low bias




Detecting overfitting

- Overfitting effect
 - Do better on training data than on future data
 - Need to choose the “right” complexity
- One solution: “Hold-out” data
- Separate our data into two sets
 - Training
 - Test
- Learn only on training data
- Use test data to estimate generalization quality
 - Model selection
- All good competitions use this formulation
 - Often multiple splits: one by judges, then another by you

Model selection


- Which of these models fits the data best?
 - $p=0$ (constant); $p=1$ (linear); $p=3$ (cubic); ...
- Or, should we use KNN? Other methods?
- Model selection problem
 - Can't use training data to decide (esp. if models are nested!)
- Want to estimate $\mathbb{E}_{(x,y)}[J(y, \hat{y}(x; D))]$

J = loss function (MSE)
 D = training data set



Hold-out method

- Validation data
 - “Hold out” some data for evaluation (e.g., 70/30 split)
 - Train only on the remainder
- Some problems, if we have few data:
 - Few data in hold-out: noisy estimate of the error
 - More hold-out data leaves less for training!




Training
data

Validation
data

$x^{(i)}$	$y^{(i)}$
88	79
32	-2
27	30
68	73
7	-16
20	43
53	77
17	16
87	94

Cross-validation method

- K-fold cross-validation
 - Divide data into K disjoint sets
 - Hold out one set ($= M / K$ data) for evaluation
 - Train on the others ($= M * (K-1) / K$ data)



Split 1:
MSE = 331.8

Split 2:
MSE = 361.2

Split 3:
MSE = 669.8


➔ 3-Fold X-Val MSE
= 464.1

Training
data
Validation
data


$x^{(i)}$	$y^{(i)}$
88	79
32	-2
27	30
68	73
7	-16
20	43
53	77
17	16
87	94

Cross-validation method


- K-fold cross-validation
 - Divide data into K disjoint sets
 - Hold out one set ($= M / K$ data) for evaluation
 - Train on the others ($= M * (K-1) / K$ data)



Split 1:
MSE = 280.5



Split 2:
MSE = 3081.3



Split 3:
MSE = 1640.1



3-Fold X-Val MSE
= 1667.3

Training
data
Validation
data

$x^{(i)}$	$y^{(i)}$
88	79
32	-2
27	30
68	73
7	-16
20	43
53	77
17	16
87	94

Cross-validation

- Advantages:
 - Lets us use more (M) validation data
(= less noisy estimate of test performance)
- Disadvantages:
 - More work
 - Trains K models instead of just one
 - Doesn't evaluate any *particular* predictor
 - Evaluates K different models & averages
 - Scores *hyperparameters / procedure*, not an actual, specific predictor!
- Also: still estimating error for $M' < M$ data...

Learning curves

- Plot performance as a function of training size

- Assess impact of fewer data on performance

Ex: $MSE_0 - MSE$ (regression)


or $1 - \text{Err}$ (classification)

- Few data

- More data significantly improve performance

- “Enough” data


- Performance saturates




- If slope is high, decreasing m (for validation / cross-validation) might have a big impact...

Leave-one-out cross-validation

- When $K=M$ (# of data), we get
 - Train on all data except one
 - Evaluate on the left-out data
 - Repeat M times (each data point held out once) and average



MSE = ...



MSE = ...

⋮



LOO X-Val MSE
= ...

Training
data
Validation
data

$x^{(i)}$	$y^{(i)}$
88	79
32	-2
27	30
68	73
7	-16
20	43
53	77
17	16
87	94

Cross-validation Issues

- Need to balance:
 - Computational burden (multiple trainings)
 - Accuracy of estimated performance / error
- Single hold-out set:
 - Estimates performance with $M' < M$ data (important? learning curve?)
 - Need enough data to trust performance estimate
 - Estimates performance of a particular, trained learner
- K-fold cross-validation
 - K times as much work, computationally
 - Better estimates, still of performance with $M' < M$ data
- Leave-one-out cross-validation
 - M times as much work, computationally
 - $M' = M-1$, but overall error estimate may have high variance

Linear Regression

Linear Regression via Least Squares

Gradient Descent Algorithms

Direct Minimization of Squared Error


Regression with Non-linear Features

Bias, Variance, & Validation

Regularized Linear Regression



What to do about under/overfitting?

- Ways to increase complexity?
 - Add features, parameters
 - We'll see more...
- Ways to decrease complexity?
 - Remove features (“feature selection”)
 - “Fail to fully memorize data”
 - Partial training
 - Regularization




Linear regression

- Linear model, two data
- Quadratic model, two data?
 - Infinitely many settings with zero error
 - How to choose among them?
- Higher order coefficients = 0?
 - Uses knowledge of where features came from...
- Could choose e.g. minimum magnitude:
$$\min \underline{\theta} \underline{\theta}^T \quad s.t. \quad J(\underline{\theta}) = 0$$
- A type of *bias*: tells us which models to prefer




Sinusoid plus noise

- Fit polynomials:




Degree 0
(constant)




Degree 1
(linear)

Sinusoid plus noise

- Fit polynomials:




Degree 3
(cubic)



Degree 9

Estimated Polynomial Coefficients

	$D=0$	$D=1$	$D=3$	$D=9$
Estimated Regression Coefficients θ	-0.03	0.75	-0.20	13.63
		-1.72	12.52	-720.38
			-40.89	12766.25
			30.43	-108900.61
				521988.15
				-1502606.92
				2648136.86
				-2794214.56
				1619195.57
				-395947.22



Regularization

- Can modify our cost function J to add “preference” for certain parameter values

$$J(\underline{\theta}) = \frac{1}{2}(\underline{y} - \underline{\theta} \underline{X}^T) \cdot (\underline{y} - \underline{\theta} \underline{X}^T)^T + \alpha \underline{\theta} \underline{\theta}^T$$

- New solution (derive the same way)

$$\underline{\theta} = \underline{y} \underline{X} (\underline{X}^T \underline{X} + \alpha \underline{I})^{-1}$$

- Problem is now well-posed for any degree

L_2 penalty:


“Ridge regression”

$$\underline{\theta} \underline{\theta}^T = \sum_i \theta_i^2$$


- Notes:

- “Shrinks” the parameters toward zero
- Alpha large: we prefer small theta to small MSE
- Regularization term is independent of the data: paying more attention reduces our model variance


Regression: Regularization



No regularization




Moderate regularization



Strong regularization


Estimated Polynomial Coefficients

	α zero	α medium	α big
<i>Estimated Regression Coefficients θ</i>	12.48	0.07	0.14
	-657.38	6.61	-0.19
	11636.48	-12.83	-0.16
	-99177.31	-9.81	-0.11
	475225.43	4.19	-0.07
	-1368006.94	11.12	-0.04
	2411278.80	9.99	0.04
	-2544758.97	4.31	0.04
	1474905.99	-2.87	0.04
	-360722.27	-9.76	0.04




Different regularization functions


- More generally, for the L_p regularizer: $\left(\sum_i |\theta_i|^p \right)^{\frac{1}{p}}$




$p=0.5$



$p=1$



$p=2$



$p=4$



Penalize small non-zero parameters more
Prefer some parameters exactly zero;
a few big parameters are OK



Penalize small non-zero parameters less,
but big parameter values a lot
Prefer lots of small weights, no big weights

“Equivalent” formulations

- Many ways to impose the same type of regularization
 - Ex: L2, quadratic penalty on parameter magnitude

$$J(\underline{\theta}) = \frac{1}{2}(\underline{y} - \underline{\theta} \underline{X}^T) \cdot (\underline{y} - \underline{\theta} \underline{X}^T)^T + \alpha \underline{\theta} \underline{\theta}^T$$

Scale by data size:

$$\frac{1}{2m}(y - \theta X^T) \cdot (y - \theta X^T)^T + \tilde{\alpha} \theta \theta^T \quad \tilde{\alpha} = \alpha/m$$

Constraint form:


$$\min_{\theta} (y - \theta X^T) \cdot (y - \theta X^T)^T \quad \text{s.t.} \quad \theta \theta^T \leq R \quad (\text{Lagrangian})$$

Magnitude vs Squared Magnitude:

$$\frac{1}{2}(y - \theta X^T) \cdot (y - \theta X^T)^T + \gamma (\theta \theta^T)^{\frac{1}{2}} \quad \gamma = \alpha \|\theta^*\|$$


Regularization: L_2 vs L_1

- Estimate balances data term & regularization term





Regularization: L_2 vs L_1

- Estimate balances data term & regularization term
- Lasso tends to generate sparser solutions than a quadratic regularizer.



Gradient-Based Optimization

- L_2 makes (all) coefficients smaller
- L_1 makes (some) coefficients exactly zero: *feature selection*



Objective Function: $f(\theta_i) = |\theta_i|^p$

Negative Gradient: $-f'(\theta_i)$

(Informal intuition: Gradient of L_1 objective not defined at zero)